

Patmos: A Time-Predictable Dual-Issue Microprocessor

Technical Report

I. INTRODUCTION

Real-time systems need a time-predictable execution platform so that the worst-case execution time (WCET) can be statically estimated. It has been argued that we have to rethink computer architecture for real-time systems instead of trying to catch up with new processors in the WCET analysis tools [2], [1].

We present the time-predictable processor Patmos as one approach to attack the complexity issue of WCET analysis. Patmos is a static scheduled, dual-issue RISC processor that is optimized for real-time systems.

II. THE ARCHITECTURE OF PATMOS

A. Pipeline

Figure 1 shows an overview of Patmos' pipeline. The pipeline consist of 5 stages: (1) instruction fetch (IF), (2) decode and register read (DR), (3) execute (EX), (4) memory access, and (5) register write back (WB).

Some instructions define additional pipeline stages. Multiplication instructions are executed, starting from the EX stage, in a parallel pipeline with fixed-length (see the instruction definition). The respective stages are referred to by EX_1, \dots, EX_n .

B. Register Files

The register files available in Patmos are depicted by Figure 2. In short, Patmos offers:

- 32, 32-bit general-purpose registers (R) : r_0, \dots, r_{31}
 r_0 is read-only, set to zero (0).
- 16, 32-bit special-purpose registers (S): s_0, \dots, s_{15}
- 8, single-bit predicate registers (P): p_0, \dots, p_7 ,
 p_0 is read-only, set to true (1).

All register reads to the R, S, and P register files are executed in the DR stage. Register writes to R are performed in the MW stage, while S and P are written immediately in the EX stage.

Concurrently writing and reading the same register in the same cycle will, for the read, yield the value that is about to be written.

When writing concurrently to the same register, i.e., the two instructions of the current bundle have the same destination register, the value of the second slot is taken, unless the predicate of that instruction evaluates to false (0).

The predicate registers are usually encoded as 4-bit operands, where the most significant bit indicates that the value read from the register file should be inverted before it is used. For operands that are written, this additional bit is omitted.

The special-purpose registers of S allow access to some dedicated registers:


- The lower 8 bits of s_0 can be used to save/restore *all* predicate registers at once. The other bits of that register are currently reserved, but not used.
- s_1 can also be accessed through the name s_m and represents the result of a decoupled load operation. The value is already sign-/zero-extended according to the load instruction. This register is read-only.
- s_2 and s_3 can also be accessed through the names s_l and s_h and represent the lower and upper 32-bits a multiplication. These registers are read-only.
- s_5 Represents the register pointing to the top of the saved stack content in the main memory.
- s_6 can also be accessed through the name s_t and represents a pointer to the top-most element of the content of the stack cache spilled to main memory. This register is read-only.

III. BUNDLE FORMATS



All Patmos instructions are 32 bits wide and are structured according to one of the instruction formats defined in the following section. Up to two instructions can be combined to form an instruction bundle; Patmos bundles are thus either 32 or 64 bits wide. The bundles sizes are recognized by the value of the most significant bit, where 0 indicates a short, 32-bit bundle and 1 a long, 64-bit bundle.

The following figures illustrate these two bundle variants:

- 32-bit bundle format



- 64-bit bundle format



IV. INSTRUCTION FORMATS

This section gives an overview of all instruction formats defined in the Patmos ISA. Individual instructions of the various formats are defined in the next section. Gray fields indicate bits whose function is determined by a sub-class of the instruction format. Black fields are not used.

- ALUi – Arithmetic Immediate

- ALUI – Long Immediate

- ALU – Arithmetic

- ALUr – Register

- ALUu – Unary

- ALUm - Multiply

- ALUc – Compare

- ALUp - Predicate

- Unused

- SPC – Special

- SPC_w - Wait

- SPCt – Move To Special

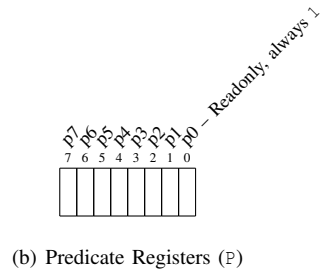
- SPCf – Move From Special

- Unused

- LDT – Load Typed

31 30 29 28 27 26 25 24 23 22 21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0
r0 (zero, read-only)
r1 (result, scratch)
r2 (result 64-bit, scratch)
r3 (argument 1, scratch)
r4 (argument 2, scratch)
r5 (argument 3, scratch)
r6 (argument 4, scratch)
r7 (argument 5, scratch)
r8 (argument 6, scratch)
r9 (scratch)
r10 (scratch)
r11 (scratch)
r12 (scratch)
r13 (scratch)
r14 (scratch)
r15 (scratch)
r16 (scratch)
r17 (scratch)
r18 (scratch)
r19 (scratch)
r20 (saved)
r21 (saved)
r22 (saved)
r23 (saved)
r24 (saved)
r25 (saved)
r26 (saved)
r27 (temp. register, saved)
r28 (frame pointer, saved)
r29 (stack pointer, saved)
r30 (function base, saved)
r31 (function offset, saved)

(a) General-Purpose Registers (R)



31 30 29 28 27 26 25 24 23 22 21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0			
reserved		p7 ... p0	s0
sm (read-only)			s1
sl (read-only)			s2
sh (read-only)			s3
s4			s6
s5			
st (read-only)			
s7			
s8			
s9			
s10			
s11			
s12			
s13			
s14			
s15			

(c) Special-Purpose Registers (S)

Fig. 2. General-purpose register files, predicate registers, and special-purpose registers of Patmos.

x	Pred	01000	Pd	Ps1	Ps2	100	Func
---	------	-------	----	-----	-----	-----	------

IF –
 DR Read register operand *Pred* and internal stack-cache registers head.
 EX Check predicate *Pred*.
 Account for head – tail < 0, update *st*.
 Update stack-cache register head.
 MW –

Note

Stack control instructions can only be issued on the first position within a bundle.

It is permissible to use several reserve, ensure, and free operations within the same function.

M. Call and Branch

Applies to CLFb and CLFi format only. Transfer control to another function or perform function-local branches. *br* performs a function-local branch within the method cache. *call* performs a function call, storing the program counter (or function offset) of the instruction to be fetched after returning in *r31*. The function base of the caller is not stored implicitly (see Section IX-C). *brcf* behaves like *call*, except that it does not write return information to *r31*.

call and *brcf* may cause a cache miss and a subsequent cache refill to load the target code; they expect the size of the code block fetched to the cache in number of bytes at <base>-4. *br* is assumed to be a cache hit.

Immediate call and branch instructions interpret the operand as *unsigned* for function calls, and as *signed* for PC-relative branches (*br*, *brcf*). The target address of PC-relative branches is computed relative to the address of the branch instruction. All immediate values are interpreted in *word size*.

Indirect call and branch instructions interpret the operand as *unsigned* absolute addresses in *byte size*.

The **link/return information** provided by *call* in *r31* should only be passed to *ret*. The unit and addressing mode (absolute or function relative) of the returned value is implementation dependent (see description of *ret*).

The following table gives an overview of the addressing modes of the available call and branch instructions.

Instruction	Immediate	Indirect	Cache fill	Link
<i>call</i>	absolute	absolute	yes	yes
<i>br</i>	PC relative	absolute	no	no
<i>brcf</i>	PC relative	absolute	yes	no

Branch and call instructions are effectively executed in the EX stage. The instructions fetched in the meantime are *not* aborted. This corresponds to a branch delay of 2 instructions that has to be respected by the compiler or assembly programmer. If no other instructions are available, two single-cycle NOPs can be used to stall the processor explicitly.

More details on the organization of the method cache is given in Section VIII.

- CLFb – Call / Branch

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
x	Pred	110	Op	Immediate																											

Op	Name	Semantics
00	<i>call</i>	function call (absolute, with cache fill)
01	<i>br</i>	local branch (PC relative, always hit)
10	<i>brcf</i>	local branch (PC relative, with cache fill)
11	—	unused

Behavior – call, brcf, and system call

IF –
 DR Read register operand *Pred*.
 EX Check predicate *Pred*.
 Store link information into *R31*. Method base is not stored to a visible register (this must be done by the caller, by convention using *R30*).
 Check method cache.
 Compute cache-relative program counter.
 If needed, fill method cache and stall.
 Update program counter.
 MW –

Behavior – branch within cache

IF –
 DR Read register operand *Pred*.
 EX Check predicate *Pred*.
 Assert on method cache.
 Compute new, cache-relative program counter value. Update program counter.
 MW –

Implementation Note

The method cache keeps track of the base address of the current function, i.e., the target/base address of the last *call*, *brcf* or *ret* instruction. *call* calculates the return offset as *nextPC_{EX}* – *base_{MC}*. However, the application code must not rely on this.

Note

All branch/call instructions can only be issued on the first position within a bundle.

- CLFi – Call / Branch Indirect

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
x	Pred			111		00										Rs1												Op			

Op	Name	Semantics
0000	<i>call</i>	function call (indirect, with cache fill)
0001	<i>br</i>	local branch (indirect, always hit)
0010	<i>brcf</i>	local branch (indirect, with cache fill)
0011	—	unused
...
1111	—	unused

Behavior – call, branch, and system call

IF –
 DR Read register operand *Pred* and *Rs1*.

EX By-pass value for `Rs1`.
 Check predicate `Pred`.
 Store link information into `R31`. Method base is not stored to a visible register (this must be done by the caller, by convention using `R30`).
 Check method cache.
 Compute cache-relative program counter.
 If needed, fill method cache and stall.
 Update program counter.

MW —

Behavior – branch within cache

IF	–
DR	Read register operand <code>Pred</code> and <code>Rs1</code> .
EX	By-pass value for <code>Rs1</code> . Check predicate <code>Pred</code> . Assert on method cache. Compute new, cache-relative program counter value. Update program counter.
MW	–

Note

All branch/call instructions can only be issued on the first position within a bundle.
See also CLFi format notes.

N. Return

Applies to CLFr format only. Transfer control to the function specified by function base and offset. `ret` may cause a cache miss and a subsequent cache refill to load the target code.

- CLFr – Return

x	Pred	111	10		Rb	Ro		Op
---	------	-----	----	--	----	----	--	----

Op	Name	Semantics
0000	ret	Return from a function (w. cache fill)
0001	—	unused
...
1111	—	unused

Behavior

IF	–
DR	Read register operand <code>Pred</code> , <code>Rb</code> and <code>Ro</code> .
EX	Check predicate <code>Pred</code> . Check method cache. Compute program counter value. If needed, fill method cache and stall. Update program counter.
MW	–

The return function base address `Rb` is an absolute address in bytes. The return function base will typically be provided by the caller in `r30` (see Section IX-C). The return function

offset `Ro` is provided by the `call` instruction in `r31`. The unit and the addressing mode of the function offset is hardware implementation dependent.

VI. DUAL ISSUE INSTRUCTIONS

Not all instructions can be executed in both pipelines. In general, the first pipeline implements all instructions, the second pipeline only a subset. All memory operations are only executed in the first pipeline.

What other instructions can be executed in both pipelines is still open for discussion and evaluation with benchmarks. A minimal approach, as first step for the hardware implementation, is to have only ALU instructions available in the second pipelines (excluding predicate manipulation instructions).

VII. STACK CACHE

The stack cache of Patmos essentially consists of a fast, small, local memory `head` and `tail` pointers into the local memory, and a top-of-stack pointer into the global memory. The structure has some similarities with a ring buffer, reserving and freeing space on the stack moves the `head` pointer, spilling and filling moves the `tail`.

As with regular ring buffers, when the size of the stack cache is not sufficient in order to reserve additional space requested, it needs to spill some data so far kept in the stack cache to the global memory, i.e., whenever $\text{head} - \text{tail} > \text{stack cache size}$. A major difference, however, is that freeing space does *not* imply the reloading of data from the global memory. When a free operation frees all stack space currently held in the cache (or more), the special register `st` is accordingly incremented.


The stack cache is organized in blocks of fixed size, e.g. 32 bytes. All spill and fill operations are performed on the block level, while reserve, free and ensure operations are in words.

Addresses for load and store operations from/to the stack cache are relative to the head pointer.

The base address for fill and spill operations of the stack cache is kept in special registers `st`.

The organization of the stack cache implies some limitations:

- The maximum size of stack data accessible at any moment is limited to the size of the cache. The stack frame can be split, such that at any moment only a subset of the entire stack frame has to be resident in the stack cache, or a *shadow* stack frame in global memory can be allocated.
- When passing pointers to data on the stack cache to other functions it has to be ensured that: (1) the data will be available in the cache, (2) the pointer is only used with load and store operations of the stack cache, and (3) the relative displacement due to reserve and free operations on the stack is known. Alternatively, aliased stack data can be kept on a *shadow* stack in the global memory without restrictions.
- The stack control operations only allow allocating constant-sized junks. Computed array sizes (C 90) and `alloca` with a computed allocation size have to be realized using a *shadow* stack in global memory.



The register usage conventions of the previous section indicate which registers are preserved across function calls.

The first 6 arguments of integral data type are passed in registers, where 64-bit integer and floating point types occupy two registers. All other arguments are passed on the *shadow* stack via the global memory.

When the return function base `r30` and the return offset `r31` needs to be saved to the stack, they have to be saved as the first elements of the function's stack frame, i.e., right after the stack frame of the calling function. Note that in contrast to `br` and `brcf` the return offset refers to the next instruction after the *delay slot* of the corresponding `call` and can be implementation dependent (cf. the description of the `call` and `ret` instructions).

D. Sub-Functions

A function can be split into several sub-functions. The program is only allowed to use `br` to jump within the same sub-function. To enter a different sub-function, `brcf` must be used. It can only be used to jump to the first instruction of a sub-function.

In contrast to `call`, `brcf` does not provide link information. Executing `ret` in a sub-function will therefore return to the last `call`, not to the last `brcf`. Function offsets however are relative to the *sub-function* base, not to the surrounding function. The function base register `r30` must therefore be set to the base address of the current *sub-function* for calls inside sub-functions.

A sub-function must be aligned and must be prefixed with a word containing the size of the sub-function, like for a regular function. If a function is split into sub-functions, the first sub-function must also be prefixed with the size of the first sub-function, not with the size of the whole function.

There are no calling conventions for jumps between sub-functions, for the compiler this behaves just like a regular jump, except that the base register `r30` must be updated if the sub-function contains calls.

E. Stack Layout

All stack data in the global memory, either managed by the stack cache or using a frame/stack pointer, grows from top-to-bottom. The use of a frame pointer is optional.

Unwinding of the call stack is done on the stack-cache managed stack frame, following the conventions declared in the previous subsection on function calls.

F. Instruction Usage Conventions

To simplify the reconstruction of the program's control flow from binary code, the use of multi-cycle NOP instructions within branch delay slots should be avoided.

X. ASSEMBLY FORMAT

A VLIW instruction consists of one or two operations that are issued in the first or both pipelines. Each operation is predicated, the predicate register is specified before the operation in parentheses (). If the predicate register is prefixed

by a !, its negation is considered. If omitted, it defaults to (p0), i.e. always true.

A double semi-colon ;; or a newline denotes the end of an instruction. If an instruction contains two operations, the operations must be separated by a single semi-colon or a single-semicolon followed by either a newline or a comment. Note that since newline is an instruction separator, the operation separator must always appear on the same line as the operation for the first slot. Labels that are prefixed by .L are local labels.

All register names must be prefixed by \$. We use destination before source in the instructions, between destination and source a = character must be used instead of a comma. Immediate values are not prefixed for decimal notation, the usual 0 and 0x formats are accepted for octal and hexadecimal immediates. Comments start with the hash symbol # and are considered to the end of the line. For memory operations, the syntax is [\$register + offset]. Register or offset can be omitted, in that case the zero register r0 or an offset of 0 is used.

Example:

```
# add 42 to contents of r2
# and store result in r1 (first slot)
add    $r1 = $r2, 42;
# if r3 equals 50, set p1 to true
cmpeq  $p1, $r3, 50
# if p1 is true, jump to label_1
($p1) br label_1 ;; nop 3    # then wait 3 cycles
# Load the address of a symbol into r2
li     $r2 = .L.str2 ;;
# perform a memory store and a pred op
swc    [$r31 + 2] = $r3 ; or $p1 = !$p2, $p3
...
label_1:
...
```

A. Inline Assembly

Inline assembly syntax is similar to GCC inline assembly. It uses %0, %1, ... as placeholders for operands. Accepted register constraints are: r or R for any general purpose register, or {<registername>} to use a specific register.

Example:

```
int i, j, k;
asm("mov  $r31 = %1 # copy i into r31\n\t"
    "add  %0 = $r5, %2"
    : "=r" (j)
    : "r" (i), "{r10}" (k));
```

REFERENCES

- [1] S. A. Edwards and E. A. Lee. The case for the precision timed (PRET) machine. In *DAC '07: Proceedings of the 44th annual conference on Design automation*, pages 264–265, New York, NY, USA, 2007. ACM.
- [2] M. Schoeberl. Time-predictable computer architecture. *EURASIP Journal on Embedded Systems*, vol. 2009, Article ID 758480:17 pages, 2009.